

Classifying albums based on dynamic range values

Sander de Jong
s2096943

January 14, 2020

1 Introduction

This project tries to classify albums based on the loudness of the songs on the album. The loudness is measured by the dynamic range values of the songs. The reason why dynamic range values could be a good classifier is found in a phenomenon called the "Loudness War". In order to make a song stand out from all others, producers and mastering engineers use all kinds of tricks to make songs sound louder when they are played at similar volume levels. This is an old phenomenon that originated in the 40s of the last century when 7-inch singles were cut to sound louder when played in jukeboxes or when songs were proposed to A&R managers of record labels[1]. When a song is louder than the competition, it makes it stand out and research has shown that humans think louder songs sound better[2].

Although it is an old phenomenon, the options to make vinyl records sound louder are limited. Excessive compression applied to vinyl records render the medium unplayable, because the needle pops out of the groove[3]. With the introduction of the compact disc (CD), these limitations were not an issue anymore and digital signal processing made it easier to apply compression to music. This made the issue more prominent, with the peak of the loudness war in the 00s. This might also be one of the reasons for the resurgence of vinyl in recent years. The perception that vinyl sounds better might have been caused by differences in mastering rather than differences in media used.

It became a topic of discussion following the release of Metallica's Death Magnetic album, which was heavily compressed[1]. The album was featured in the rhythm game Guitar Hero in which players try to play along with the record. The game used the unmastered multi-tracks to provide the sound of the different instruments that could be played in game. The files were ripped from the game and mixed together, resulting in a version of the album with far less compression. The excessive compression used and the fact that it concerned a popular band made the issue apparent to a larger audience. The downward

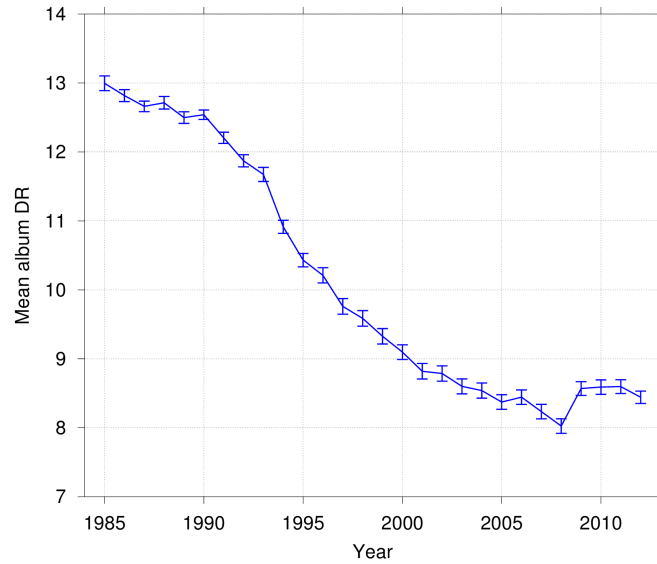
trend seems to have stopped after that, seeing wider dynamic range in records from the last ten years.

Because the phenomenon became more apparent after the introduction of the CD and was most prevalent in the 90s and 00s, it is interesting to see the trend from 1985 onwards. Inspiration for this research was drawn from an article written in 2013[4] in which Jasper van Dorp visualised this trend based on 14680 values. The goal of this project is to see whether these values can be used as a classifier to predict the year an album was released in. Further inspection of the database raised the suspicion that the values of albums in different genres might vary enough for them to be reliable classifiers.

Therefore this project consists of two experiments with the following hypotheses:

- H1: The database of dynamic range values can be used as classifier to determine the year an album was released in based on its dynamic range value.
- H2: The dynamic range values of albums in different genres vary enough to use it as a classifier to determine the genre an album with a specific value belongs to.

Figure 1: The development of dynamic range values from 1985 to 2014 as measured by Jasper van Dorp



2 Methods

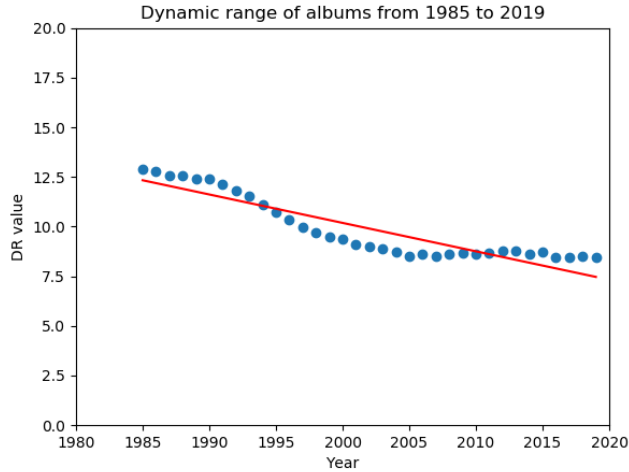
2.1 Algorithm used to calculate dynamic range

The algorithm that is used to calculate dynamic range values to be submitted to the database is DROffline[5]. It calculates the peak-to-loudness ratio or crest factor. It evaluates this crest factor throughout the whole song, resulting in a single value for the song. When the algorithm is presented with an album of songs, it outputs the mean dynamic range values of the songs on the album along with the individual song values.

2.2 Dataset

Hi-fi enthusiasts have collaboratively maintained a complete database of dynamic range values consisting of over 140000 albums (<http://dr.loudness-war.info>). Even though the database has thrived due to submissions by volunteers, the maintainer of the website doesn't provide an API to use the information and is not reachable through email or social media. Therefore the values had to be obtained by building a custom web scraper that saves the dynamic range values of all entries (for the code of the web scraper see the appendix). The data used in these experiments were scraped on 29 November 2019.

Figure 2: The development of dynamic range values from 1985 to 2019, based on the database entries



The mean DR values of the albums in the years from 1985 until 2019 are plotted in figure 2 using blue dots. It shows roughly the same curve as Jasper van Dorp's, although this database consists of a lot more values. There are some problems with the dataset that influence the results. Although over half of the changes

in values throughout the years are statistically significant (see appendix A), the spread is not that large. Moreover, the values submitted to the database are rounded to integers instead of the decimal values that result from the algorithm used to obtain the dynamic range values.

Figure 3: The number of dynamic range values present in the database for the various years

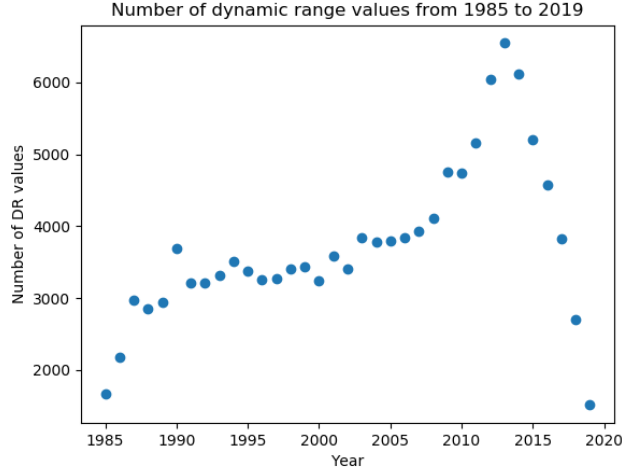


Figure 3 shows that although there are more values for recent years, all years contain enough data points for this experiment. Apart from the limitation in years used, vinyl records were excluded due to several reasons. As mentioned in the introduction, the loudness war isn't as prevalent as it is with CDs. Furthermore, the algorithm used doesn't work as well for vinyl records[6]. By only using CDs, the format is kept constant to test the other parameters.

2.2.1 Genre data

Because the dynamic range database doesn't contain information for the genres of the albums, the scraped data had to be cross-referenced with another database containing this information. After evaluating APIs of different services, Last.fm was chosen because it is quite lenient concerning rate limiting[7]. It specifies a limit of one API call per second, but the API was used with a higher rate without any problems. Another reason why this service was chosen is because it contains genre data for artists instead of individual songs. To reduce the time it takes to get the data, the decision was made to obtain genre data for artists instead of for individual albums. Although this is a simplification because artists can make albums in various genres, this limited the amount of items to be requested to 30000 instead of 140000. Even with this simplification, obtaining these 3000 artist-genre combinations took 8 hours. The genre data Last.fm provides is

submitted by users. It maintains a list of genres assigned by users and a count of how many users chose this specific genre. The most popular genre for each artist was linked to its albums.

Figure 4: An extract of the database cross-referenced with the genre data from Last.fm

12271	jazz	A cry from the heart - Live in London	15.0	1988.0
12272	jazz	Picture This	15.0	1987.0
12273	jazz	Warning	16.0	1985.0
12274	pop	Both Sides	11.0	1993.0
12275	pop	No Jacket Required	14.0	1985.0
12276	pop	Twenty Four Seven	9.0	1999.0
12277	pop	Foreign Affair	13.0	1989.0
12278	pop	Wildest Dreams	7.0	1996.0
12279	pop	Simply The Best	10.0	1991.0
12280	pop	Live	6.0	2009.0
12281	pop	The Bridge	9.0	1995.0
12282	jazz	Matt Bianco	14.0	1986.0
12283	pop	New Arrival	6.0	2003.0
12284	pop	Pop 'Til You Drop!	6.0	2002.0
12285	pop	Teen Spirit	6.0	2001.0
12286	pop	The Abba Generation	6.0	1999.0
12287	pop	(What's the Story) Morning Glory?	6.0	1996.0
12288	jazz	New Moon Daughter	12.0	1995.0
12289	pop	Mi Plan	7.0	2009.0
12290	jazz	Unforgettable Glenn Miller	12.0	1985.0

2.3 Experiments

2.3.1 Experiment I: Determining the date of an album based on dynamic range values

The values of the dataset are used as a classifier to find the year an album was released. This is done by creating a regression line using the mean dynamic range value per year. The regression line can be seen in red in figure 2. The means are used as training data, while the values of the individual albums are used as test data.

Furthermore, a feed-forward neural network model was used to predict the year based on the available dynamic range values. The network consists of two dense layers with ReLu activation and a last dense layer with softplus activation. Several activation functions were tested, but the outcomes were similar. For the neural network, a different split in training and validation data is used. The means that were used by the linear regression are not needed here and therefore the individual values are separated in training and validation data. The full code for the linear regression and neural network can be found in the appendix.

2.3.2 Experiment II: Determining the genre of an album based on its dynamic range values

The genre data from Last.fm contains quite specific genres. To limit the number of genres compared, several subgenres are combined to create logical categories. For example, all subgenres with "Rock" in the genre name are combined into one category. The same for pop, blues, metal etc. Several of these categories are

combined to see whether the dynamic range values can be used to distinguish them from each other.

The resulting input data is split into training and test data. Subsequently, several machine learning classifiers are used to see which method works best for this experiment. Finally, a feed-forward neural network is used to see whether this gives better results than the built-in classifiers. The network consists of three layers, two dense layers with 512 nodes and a final dense layer with nodes equal to the amount of genres that are evaluated. The intermediate layers use ReLu activation and the output layer uses softmax activation, because this works best for multi-class classification problems. Different amount of nodes and different activation functions were tested, but these didn't make an impact or even made the performance worse.

The neural network is trained with two different datasets. One with just the genre and dynamic range data and one test with genre, dynamic range and year data.

3 Results

3.1 Experiment I

The results of the linear regression show a Root Mean Squared Error of 15 years, which is not very useful as classifier. This is not that surprising considering the aforementioned problems with the dataset and the fact that linear regression can result in values outside of the specified range. It can even predict years in the future, therefore the large error rate is not surprising.

Results of testing regression:

Mean Absolute Error: 11.25

Mean Squared Error: 230.9

Root Mean Squared Error: 15.20

The neural network predicts the year with an accuracy of 6% which is higher than the chance value of 3%, but is not very useful. The flaws with the dataset are too big to get impressive results using this method. Tweaking the network didn't change the accuracy. Because the results of experiment II were more promising, the focus shifted towards that experiment.

Results of the feed-forward neural network:

Training accuracy: 0.0616

Validation accuracy: 0.0631

3.2 Experiment II

The various classifiers work very well to predict the genre of an album based on its dynamic range (see table 1). The classifiers give similar results and the neural network doesn't bring improvements over the classifiers. This is to be expected since the classifiers are built specifically for these kinds of problems. Tweaks to the neural network didn't provide better results and more advanced types of neural networks seem like an overkill for this problem. The prediction tree method gives the best scores, although the differences are small. The SVM takes significantly longer to calculate the result and performs worse, this is due to the fact that SVMs are best used for problems with multiple input variables. The addition of year data gives far worse results compared to just using the genre and dynamic range values, therefore these results have been omitted and this part of the experiment has not been pursued further. It shows once again that the combination of year and dynamic range values is not very useful as classifier.

Genre data does provide promising results. Genres of completely different music work best. Especially classical music can be discriminated from the other genres with near-perfect accuracy. When genres of more similar music are classified together, the performance is worse but still quite impressive. Rock and metal can sound quite similar, but a combination of the two still scores 68%.

Genres	KNN	Prediction Tree	Naive Bayes	SVM	Feed-forward NN
Pop, Rock, Jazz	0.6188	0.7048	0.6961	0.6961	Train: 0.6953 Test: 0.7005
Pop, Rock, Metal, Jazz	0.5020	0.5354	0.5293	0.5072	Train: 0.5376 Test: 0.5361
Pop, Rock, Metal, Jazz, Classical	0.5371	0.5411	0.5336	0.5085	Train: 0.5304 Test: 0.5444
Pop, Rock, Metal, Jazz, Classical, Blues, Hip-hop	0.5067	0.5071	0.5055	0.4831	Train: 0.5133 Test: 0.5104
Pop, Jazz, Classical	0.6931	0.7410	0.7314	0.7317	Train: 0.7384 Test: 0.7410
Rock, Metal	0.6826	0.6826	0.6812	0.6500	Train: 0.6865 Test: 0.6829
Pop, Jazz	0.6506	0.7489	0.7450	0.7450	Train: 0.7478 Test: 0.7450
Metal, Jazz	0.7593	0.7708	0.7730	0.7730	Train: 0.7716 Test: 0.7688
Rock, Jazz	0.7269	0.8110	0.8031	0.8031	Train: 0.7965 Test: 0.8110
Rock, Classical	0.9942	0.9944	0.9941	0.9942	Train: 0.9946 Test: 0.9944
Pop, Classical	0.9779	0.9779	0.9764	0.9764	Train: 0.9718 Test: 0.9779

Table 1: The results for the different combinations of genres using various classifiers.

4 Conclusion

The initial goal of the project to date albums based on dynamic range values was not successful. This can be attributed to problems with the dataset. The website it was scraped from rounds the values to integers, making it less precise than the algorithm provides. But even if the values would be available as floats, the spread is still quite small making it hard to use it as a classifier. Furthermore, linear regression might not be the best approach to this problem because it can result in values outside of the specified range. If the experiment would be repeated with better data and other approaches, the expectation is that the results would be better but still not very impressive.

During the research of dynamic range values, it became apparent that genres have different mean values and that this might be a better classifier. The experiments showed that this is the case, classifying a combination of classical music and rock or pop music with near perfect scores. Although classical music is completely different from other genres, even more similar genres like rock and

metal can be distinguished from each other very accurately.

References

- [1] All Things Considered show. The Loudness Wars: Why Music Sounds Worse. 31 December 2009. Accessed 13 January 2020. <https://www.npr.org/2009/12/31/122114058/the-loudness-wars-why-music-sounds-worse?sc=nl&cc=mn-20100102&t=1578856508222>
- [2] Welch, D. and Fremaux, G. (2017). Why Do People Like Loud Sound? A Qualitative Study. *Int J Environ Res Public Health*. 2017 Aug; 14(8): 908. Published online 2017 Aug 11. doi:10.3390/ijerph14080908
- [3] Gross, J. (2006). Everything Louder Than Everything Else: Have the loudness wars reached their final battle? *XL Recording Studio Guide* 2006. Accessed at <https://web.archive.org/web/20061019013037/http://www.austin360.com/music/content/music/stories/xl/2006/09/28cover.html> on 13 January 2020
- [4] Van der Dorp, J. (2013). The Loudness War: Dynamic Range over the years. *TheMindGap*. Accessed at <http://www.themindgap.nl/?p=114> on 13 January 2020.
- [5] DROffline. MAAT. Accessed at <https://www.maat.digital/droffline/> on 14 January 2020
- [6] Sheperd, I (2013). Why you can't measure vinyl with the TT Meter. *Production advice*. accessed at <https://productionadvice.co.uk/tt-meter-not-for-vinyl/> on 14 January 2020.
- [7] Last FM API. Accessed at <https://www.last.fm/api/> on 14 January 2020.

Appendix A: p-values

Year	p-value	Significant
1985	0.0922	False
1986	0.0	True
1987	0.8663	False
1988	0.0001	True
1989	0.3855	False
1990	0.0	True
1991	0.0	True
1992	0.0	True
1993	0.0	True
1994	0.0	True
1995	0.0	True
1996	0.0	True
1997	0.0001	True
1998	0.0	True
1999	0.0626	False
2000	0.0	True
2001	0.0216	True
2002	0.2059	False
2003	0.0017	True
2004	0.0001	True
2005	0.0313	True
2006	0.0421	True
2007	0.0858	False
2008	0.1411	False
2009	0.465	False
2010	0.5137	False
2011	0.1187	False
2012	0.4518	False
2013	0.0008	True
2014	0.0308	True
2015	0.0	True
2016	0.8161	False
2017	0.4694	False
2018	0.9208	False

Table 2: The p-values for the comparison between subsequent years. The p-value is the outcome of a t-test between the dynamic range values of the year mentioned and the year after

Appendix B: code

```
1 from selenium import webdriver
2 from BeautifulSoup import BeautifulSoup
3 import pandas as pd
4
5 driver = webdriver.Chrome()
6 results = [[]]
7 # nr_pages should contain the amount of pages of http
8   ://dr.loudness-war.info/album/list/
9 # While scraping the data this value was 1391
10 nr_pages = 1391
11 def add_page(url):
12     driver.get(url)
13     content = driver.page_source
14
15     # Use the BeautifulSoup package to find the right
16     # column and save its values in results
17     bs=BeautifulSoup(content)
18     div = bs.find('div', {'class': 'table-responsive-sm'})
19
20     if div is not None:
21         rows = div.findAll('tr')
22         #print(rows)
23         for row in rows:
24             tds=row.findAll('td')
25             result = []
26             for td in tds:
27                 if td:
28                     data = td.text.strip()
29                     if data:
30                         result.append(data)
31             results.append(result)
32
33 for i in range(1,nr_pages+1):
34     add_page("http://dr.loudness-war.info/album/list/" +
35             str(i))
36 # Create a dataframe from results and save it to disc.
37 df = pd.DataFrame(results, columns=["Artist", "Album",
38     "Year", "avg DR", "min DR", "max DR", "codec", "
39     source"])
40 print(df)
41 df.to_csv('dr_data.csv', encoding="utf-8")
```

Listing 1: scrapedr.py

```

1 import numpy as np
2 from sklearn.linear_model import LinearRegression
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import discogs_client
6 import pylast
7 import os, time
8 import requests
9
10
11 # Read the dr data obtained in scrape_dr and find all
    the artists.
12 data = pd.read_csv("dr_data.csv")
13 data = data[data['Year'] >= 1985.0]
14 artists = data['Artist'].unique()
15 # A Last.fm API key is needed to use this script.
16 API_KEY =
17 URL = "http://ws.audioscrobbler.com/2.0/"
18 # Find an artist and return its most popular genre.
19 def get_genre(artist, PARAMS):
20     r = requests.get(url = URL, params = PARAMS).json()
21
22     try:
23         genre = list(list(r.values())[0].values())[0][0]['
            name']
24         return genre
25     except:
26         return "Not found"
27
28 genres = np.empty([len(artists), 2], dtype="U200")
29 for i in range(0, len(artists)):
30     artist = artists[i]
31     PARAMS = {'method': 'artist.gettoptags', 'artist' :
        artist, 'api_key':API_KEY, 'format':'json'}
32     genre = get_genre(artist, PARAMS)
33     if(genre != "Not found"):
34         genres[i] = [artist, genre]
35         np.savetxt('genres.txt', genres, fmt='%s',
            delimiter=',')

```

Listing 2: getgenres.py

```

1 import csv, pandas as pd
2
3 # Open the dr data obtained in scrape_dr.py and the
    genre data obtained in get_genre.py
4 data = pd.read_csv("dr_data.csv")

```

```

5 data = data[data['Year'] >= 1985.0]
6 data = data[data['source'] != "Vinyl"]
7
8
9 with open('genres.csv', mode='r') as infile:
10     reader = csv.reader(infile)
11     genres = {rows[0]:rows[1] for rows in reader}
12
13 temp_df = [[]]
14 # Iterate over the albums and combine all subgenres
15 # into logical categories.
16 for index, row in data.iterrows():
17     try:
18         if "pop" in genres[row['Artist']]:
19             temp_df.append(["pop", row['Album'], row['avg DR'], row['Year']])
20         if "rock" in genres[row['Artist']]:
21             temp_df.append(["rock", row['Album'], row['avg DR'], row['Year']])
22         if "metal" in genres[row['Artist']]:
23             temp_df.append(["metal", row['Album'], row['avg DR'], row['Year']])
24         if "jazz" in genres[row['Artist']]:
25             temp_df.append(["jazz", row['Album'], row['avg DR'], row['Year']])
26         if "classical" in genres[row['Artist']]:
27             temp_df.append(["classical", row['Album'], row['avg DR'], row['Year']])
28         if "hiphop" in genres[row['Artist']] or "rap" in genres[row['Artist']]:
29             temp_df.append(["rap", row['Album'], row['avg DR'], row['Year']])
30         if "blues" in genres[row['Artist']]:
31             temp_df.append(["blues", row['Album'], row['avg DR'], row['Year']])
32     except:
33         print("Artist not in database")
34 # Save the combination of album, dr and genre data
35 data = pd.DataFrame(temp_df, columns = ["Genre", "Album", "DR", "Year"])
36 data.to_csv('genre_data.csv', encoding="utf-8")

```

Listing 3: combinedata.py

```

1 import pandas
2 from keras.models import Sequential
3 from keras.layers import Dense

```

```

4 from keras.wrappers.scikit_learn import
   KerasClassifier
5 from keras.utils import np_utils
6 from sklearn.model_selection import cross_val_score
7 from sklearn.model_selection import KFold
8 from sklearn.preprocessing import LabelEncoder
9 from sklearn.pipeline import Pipeline
10 import numpy as np
11 from sklearn.neighbors import KNeighborsClassifier
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import confusion_matrix
14 from sklearn.tree import DecisionTreeClassifier
15 from sklearn.metrics import accuracy_score
16 from sklearn.naive_bayes import GaussianNB
17 from sklearn.svm import SVC
18
19 # Read the genre data from combine_data.py and extract
   the columns that are needed
20 dataframe = pandas.read_csv("genre_data", header=None)
21 dataframe = dataframe.dropna()
22 dataset = dataframe.values
23 X = dataset[:, [3]].astype(float)
24 Y = dataset[:, 1]
25 print(X)
26 print(X.shape)
27 print(Y.shape)
28
29 X_train, X_test, y_train, y_test = train_test_split(X,
   Y, random_state = 0)
30
31 # Use the K-nearest-Neighbours classifier and return
   the
32 # resulting accuracy and confusion matrix
33 knn = KNeighborsClassifier(n_neighbors = 7).fit(
   X_train, y_train)
34 accuracy = knn.score(X_test, y_test)
35 print(accuracy)
36 knn_predictions = knn.predict(X_test)
37 cm = confusion_matrix(y_test, knn_predictions)
38 print(cm)
39
40
41 # Use the Decision Tree classifier and return the
42 # resulting accuracy and confusion matrix
43 dtree_model = DecisionTreeClassifier(max_depth = 2).
   fit(X_train, y_train)

```

```

44 dtree_predictions = dtree_model.predict(X_test)
45 accuracy = accuracy_score(y_test, dtree_predictions)
46 print(accuracy)
47
48 cm = confusion_matrix(y_test, dtree_predictions)
49 print(cm)
50
51
52
53 # Use the Gaussian Naive Bayes classifier and return
    the
54 # resulting accuracy and confusion matrix
55 gnb = GaussianNB().fit(X_train, y_train)
56 gnb_predictions = gnb.predict(X_test)
57
58 accuracy = gnb.score(X_test, y_test)
59 print(accuracy)
60
61 cm = confusion_matrix(y_test, gnb_predictions)
62 print(cm)
63
64
65 # Use the Support Vector Machine classifier and return
    the
66 # resulting accuracy and confusion matrix
67 svm_model_linear = SVC(kernel = 'linear', C = 1).fit(
    X_train, y_train)
68 svm_predictions = svm_model_linear.predict(X_test)
69
70 accuracy = svm_model_linear.score(X_test, y_test)
71 print(accuracy)
72
73 cm = confusion_matrix(y_test, svm_predictions)
74 print(cm)
75
76 #Convert the potential genres outcomes using one-hot
    encoding
77 encoder = LabelEncoder()
78 encoder.fit(Y)
79 encoded_Y = encoder.transform(Y)
80 dummy_y = np_utils.to_categorical(encoded_Y)
81
82 X_train, X_test, y_train, y_test = train_test_split(X,
    dummy_y, random_state = 0)
83 def baseline_model():
84     model = Sequential()

```

```

85     model.add(Dense(512, input_dim=1, activation='relu')
86         )
87     model.add(Dense(512, activation='relu'))
88     model.add(Dense(3, activation='softmax'))
89     model.compile(loss='categorical_crossentropy',
90         optimizer='adam', metrics=['accuracy'])
91     return model
92 # Build a simple feed-forward neural network and test
93 # its performance on the genre and dr data
94 model = baseline_model()
95 model.summary()
96 model.fit(X_train, y_train, batch_size=100, nb_epoch
97     =25, validation_data=(X_test, y_test))

```

Listing 4: classifier.py

```

1  import numpy as np
2  from sklearn.linear_model import LinearRegression
3  import pandas as pd
4  import matplotlib.pyplot as plt
5  from sklearn.model_selection import train_test_split
6  from sklearn.metrics import accuracy_score
7  from sklearn import metrics
8  from sklearn.preprocessing import LabelEncoder
9  from sklearn.pipeline import Pipeline
10 from keras.utils import np_utils
11 from keras.models import Sequential
12 from keras.layers import Dense
13 from keras.wrappers.scikit_learn import
14     KerasClassifier
15 # Read the DR data that was scraped using scrape_dr.py
16 # and
17 # prepare the means as training data and the
18 # individual values
19 # as test values.
20 data = pd.read_csv("dr_data.csv")
21 data = data[data['source'] != 'Vinyl']
22 data = data[['Year', 'avg DR']]
23 data = data[data['Year'] >= 1985.0]
24 data_mean = data.groupby(['Year'])['avg DR'].mean()
25 data_mean = pd.DataFrame(data_mean, columns = ['avg DR
26     '])
27 dr_train = np.array(data_mean.iloc[:,0]).reshape((-1,
28     1))
29 year_train = np.array(data_mean.index.values).reshape
30     ((-1, 1))

```



```

26
27 dr_test = np.array(data['avg DR']).reshape(-1,1)
28 year_test = np.array(data['Year'])
29
30 # Perform the linear regression and print the outcome
31 model = LinearRegression().fit(dr_train, year_train)
32 r_sq = model.score(dr_train, year_train)
33
34 print('coefficient of determination:', r_sq)
35 print('intercept:', model.intercept_)
36 print('slope:', model.coef_)
37
38 y_pred = model.predict(dr_test)
39 df = pd.DataFrame({'Actual': year_test.flatten(), '
    Predicted': y_pred.flatten()})
40 print(df)
41
42 # Plot the means per year and the regression line
43 plt.scatter(year_train, dr_train)
44 axes = plt.gca()
45 axes.set_xlim([1980,2020])
46 axes.set_ylim([0,20])
47 axes.set_title("Dynamic range of albums from 1985 to
    2019")
48 axes.set_xlabel("Year")
49 axes.set_ylabel("DR value")
50 y_pred = model.predict(year_train)
51 plt.plot(year_train, y_pred, color='red')
52 plt.show()
53
54 print('Mean Absolute Error:', metrics.
    mean_absolute_error(year_train, y_pred))
55 print('Mean Squared Error:', metrics.
    mean_squared_error(year_train, y_pred))
56 print('Root Mean Squared Error:', np.sqrt(metrics.
    mean_squared_error(year_train, y_pred)))
57
58 # Encode the potential output years with one-hot
    encoding
59 encoder = LabelEncoder()
60 encoder.fit(year_test)
61 encoded_Y = encoder.transform(year_test)
62 dummy_y = np_utils.to_categorical(encoded_Y)
63 X_train, X_test, y_train, y_test = train_test_split(
    dr_test, dummy_y, random_state = 0)
64

```

```

65
66 def baseline_model():
67     model = Sequential()
68     model.add(Dense(512, input_dim=1, activation='relu')
69     )
69     model.add(Dense(512, activation='relu'))
70     model.add(Dense(35, activation='softmax'))
71     model.compile(loss='categorical_crossentropy',
72                   optimizer='adam', metrics=['accuracy'])
72     return model
73
74 # Build a simple feed-forward neural network and test
75   its performance on the year and dr data
76 model = baseline_model()
77 model.summary()
78 model.fit(X_train, y_train, batch_size=100, nb_epoch
79           =25, validation_data=(X_test, y_test))

```

Listing 5: predictor.py